AgentCamp

28 February 2026

10:00 AM – 03:00 PM CST

HROne Office
Noida, India.

agentcamp.city/noida

An event by the: global AI community

AgentCamp

Join the AI Conversation
on Discord

gaic.io/discord

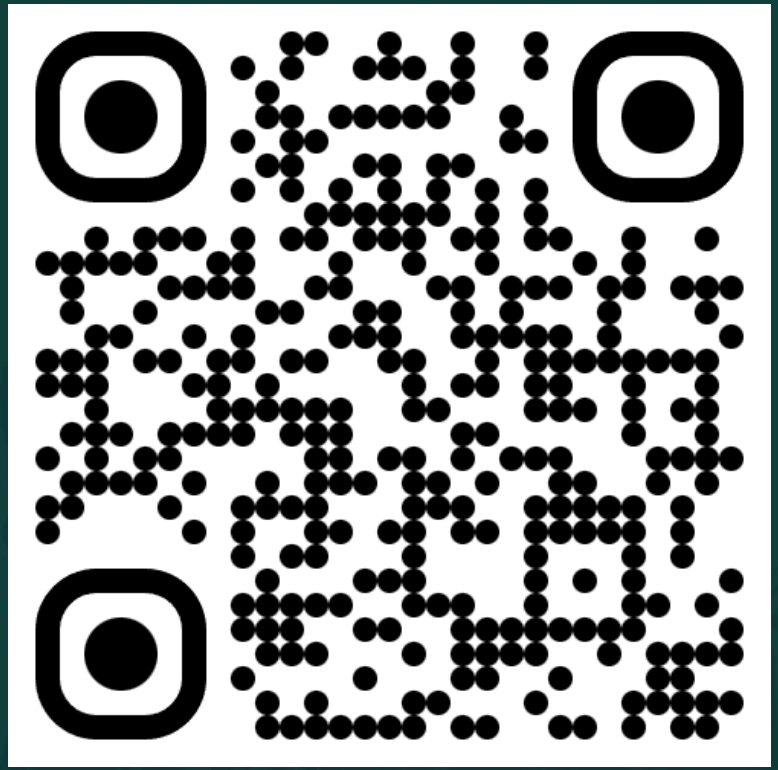

An event by the: global AI community

AgentCamp

Social Media Hashtags

#AgentCamp

#GlobalAICommunity

AgentCamp

How Smart AI Agents

Think, Act & Automate Tasks for You

Sunny Sharma

Cloud & Solutions Architect | DevOps · Observability · Automation

An event by the: global AI community

What We'll Cover

Chatbot vs. Agent

Why agents are fundamentally different

The ReAct Loop

How agents think, decide, and act

Tool Use

Calling APIs, databases, and services

RAG

How agents retrieve knowledge

MCP

The universal connector standard

Real-World Patterns

Multi-agent, reflection, safety

LLMs Are Powerful. But Limited.

Large Language Models can generate remarkable text, but they can't act on the world by themselves.

Knowledge Cutoff

Frozen at training date.
No real-time info.

No Data Access

Can't see your files,
databases, or systems.

No Actions

Can't send emails,
call APIs, or deploy.

No Memory

Every conversation
starts from scratch.

→ *This is exactly why we need AI Agents — systems that can think AND act.*

Chatbot vs. AI Agent

Chatbot

- Responds to what you say
- Reactive, one turn at a time
- Cannot take real actions
- Generates text only
- No access to external data

AI Agent

- Works toward a goal
- Proactive, multi-step
- Calls APIs & takes actions
- Uses tools & retrieves data
- Adapts based on observations

What Makes an AI Agent?

LLM
The Brain

+

Tools
The Hands

+

Memory
The Context

+

Reasoning
The Loop

= An AI system that understands goals, makes plans, uses tools, and takes real actions

The ReAct Loop

Reasoning + Acting

1 THOUGHT

The agent reasons about the current state and decides what to do next

→

2 ACTION

The agent calls a tool: API, database, search, code execution, etc.

→

3 OBSERVATION

The agent receives the tool's result and evaluates if the goal is met

↻ Repeat until the goal is achieved or a stop condition is met

Analogy: Like a GPS — drive a bit, recalculate, adjust. The agent thinks, acts, observes, adjusts.

ReAct in Action: A Real Example

User: "What's the shipping status of my last order?"

THOUGHT	"I need the user's most recent order. Let me query the orders database."
ACTION	Call orders_api.get_recent(user_id=123) → Returns Order #789
THOUGHT	"Got Order #789. Now I need its shipping status."
ACTION	Call shipping_api.status(order_id=789) → In Transit, ETA: March 2
THOUGHT	"I have all the info. Let me compose the final response."
ANSWER	"Your last order (#789) is in transit and should arrive by March 2."

Tool Use — The Agent's Hands

The LLM decides which tool to call. The framework executes it.

1 → 2 → 3 → 4

Define

Describe tools to the LLM

Decide

LLM outputs JSON choice

Execute

Framework runs the tool

Return

Result back to the LLM

Common Agent Tools

- Web Search
- Database
- Code Exec
- Email/Slack
- REST APIs
- File Search

RAG — Retrieval Augmented Generation

Give the LLM access to YOUR data at query time — no retraining needed.

1

RETRIEVE

Search your knowledge base (vector DB) for relevant documents

→

2

AUGMENT

Combine retrieved context with the user's question into one prompt

→

3

GENERATE

LLM produces an answer grounded in actual data, not guesses

RAG is not separate from agents — it is one of the most common TOOLS an agent uses.

RAG vs. Fine-Tuning: RAG is cheap, fast, updatable, cites sources. Fine-tuning is expensive, slow, permanent.

MCP — Model Context Protocol

The USB-C of AI — a universal connector for agents and tools.

Before MCP

5 AI Apps × 10 Tools = 50 Custom Connectors

- Every integration built from scratch
- Different format per provider
- Maintenance nightmare at scale

With MCP

5 Clients + 10 Servers = 15 Implementations

- Build once, connect everywhere
- Standardized JSON-RPC protocol
- Dynamic tool discovery at runtime

Now an industry standard: Adopted by OpenAI (Mar 2025), donated to Linux Foundation (Dec 2025).

MCP Architecture

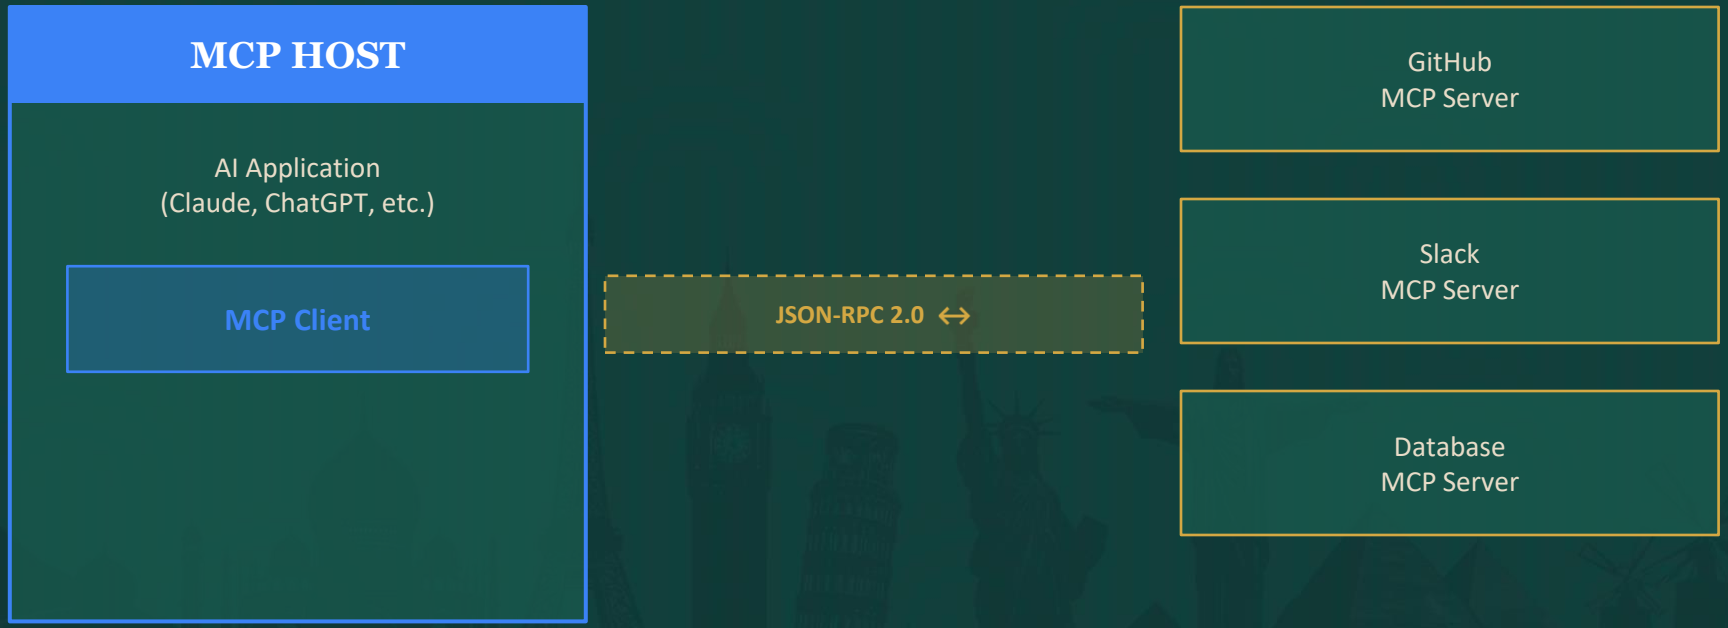

The agent never needs to know how each API works — MCP servers handle all the complexity.

Real-World Agent Patterns

Multi-Agent

Specialized agents collaborate:
researcher, writer, reviewer.
Like a human team.

Reflection

Agent reviews its own output,
catches errors, and self-corrects
before delivering.

Human-in-the-Loop

Agent pauses at critical decisions
for human approval.
Trust through oversight.

Plan & Execute

Create full plan first,
then execute step by step.
Big picture before details.

Your Learning Path

Where to start if you want to build agents:

1. **Learn LLM APIs** →
2. **Prompt Engineering** →
3. **Build a RAG System** →
4. **Add Tool Use & MCP** →
5. **Build Your First Agent**

Frameworks & Tools to Explore

- **LangChain** Most popular agent framework
- **LangGraph** Stateful agent workflows
- **CrewAI** Multi-agent collaboration
- **Ollama** Local LLM execution
- **FastMCP** Build MCP servers easily
- **Pinecone / Qdrant** Vector databases for RAG

AgentCamp

Thank You!

Questions & Discussion

Sunny Sharma

linkedin.com/in/sunnyksharma

#AgentCamp #GlobalAICommunity